

amazon_reviews_eda

March 19, 2021

```
[1]: from pyspark.sql import SQLContext, SparkSession
      from pyspark.sql.types import *
      from pyspark.sql.functions import *
      from pyspark import SparkContext, SparkConf
```

```
[2]: spark = SparkSession.builder.getOrCreate()
      sc = spark.sparkContext
```

```
[3]: filename = 'kindle_reduced_clean.csv'
      df = spark.read.csv(filename, inferSchema=True, header = True)
```

```
[4]: df.select("overall", "summary", "reviewText").show(5)
```

```
+-----+-----+-----+
|overall|          summary|      reviewText|
+-----+-----+-----+
|      5| A Very Sexy Cruise|ARC provided by a...|
|      5|A Changing Gears ...|Wild Ride by Nanc...|
|      5|We don't take kin...|Well thought out ...|
|      3|Mediocre Science ...|Being autistic, I...|
|      3| I'm losing interest|This is book four...|
+-----+-----+-----+
only showing top 5 rows
```

```
[5]: df.select([count(when(col(c).isNull(), c)).alias(c) for c in df.columns]).show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|index|asin|helpful|overall|reviewText|reviewTime|reviewerID|reviewerName|summary|
|unixReviewTime|HelpfulRecords|HasHelpful|weightedRating|
+-----+-----+-----+-----+-----+-----+-----+-----+
|    0|    0|    0|    0|    1|    0|    0|    24|
|    0|    0|    0|    0|    0|    0|    0|
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
```

```
[6]: df = df.dropna(how='any')

[7]: df=df.drop("index","reviewerName","unixReviewTime","helpful","HasHelpful")
```

```
[8]: df = df.withColumn('reviewText', translate('reviewText', '.', ''))
df = df.withColumn('reviewText', translate('reviewText', ',', ''))
df = df.withColumn('reviewText', translate('reviewText', '$', ''))
```

```
[9]: from pyspark.ml.feature import Tokenizer, StopWordsRemover

#tokenize text (make words into an array)
tokenizer = Tokenizer(inputCol='reviewText', outputCol='reviewText_token')
df_token = tokenizer.transform(df).select('*')

#remove basic words
remover = StopWordsRemover(inputCol='reviewText_token',
    ↳outputCol='reviewText_clean')
df_stop=remover.transform(df_token).select('*')
```

```
[10]: #tokenize summaries (make words into an array)
tokenizer = Tokenizer(inputCol='summary', outputCol='summary_token')
df_token = tokenizer.transform(df_stop).select('*')

#remove basic words
remover = StopWordsRemover(inputCol='summary_token', outputCol='summary_clean')
df_stop=remover.transform(df_token).select('*')
```

```
[11]: df_stop=df_stop.drop("reviewText", "summary","reviewText_token",
    ↳"summary_token")
df_stop.show(5)
```

```
+-----+-----+-----+-----+-----+-----+-----+
|      asin|overall| reviewTime|      reviewerID|HelpfulRecords|weightedRating|
reviewText_clean|      summary_clean|
+-----+-----+-----+-----+-----+-----+-----+
|B00J4S6YWC|      5|06 21, 2014| AUSBN91MCI3WM|      0.0|
5.0|[arc, provided, a...|      [sexy, cruise]|
|B00HCZUBH8|      5| 03 3, 2014|A141H51I3H4B1S|      0.5|
5.0|[wild, ride, nanc...|[changing, gears,...|
|B006RZNR3Y|      5|07 10, 2014| AP8TKDM76TROZ|      0.0|
4.0|[well, thought, s...| [take, kindly, no!]|
|B006RZNR3Y|      3| 02 1, 2014|A22GGHISKRVAOX|      0.0|
4.0|[autistic, freque...|[mediocre, scienc...|
|B00J47H8H8|      3|03 21, 2014|A19DWIC1T7127Y|      0.75|
3.0|[book, four, five...| [losing, interest]|
```

```
+-----+-----+-----+-----+-----+-----+-----+
|-----+-----+
only showing top 5 rows
```

```
[12]: display(df_stop.select("reviewText_clean"))
```

```
DataFrame[reviewText_clean: array<string>]
```

```
[13]: df_stop.printSchema()
```

```
root
|-- asin: string (nullable = true)
|-- overall: integer (nullable = true)
|-- reviewTime: string (nullable = true)
|-- reviewerID: string (nullable = true)
|-- HelpfulRecords: double (nullable = true)
|-- weightedRating: double (nullable = true)
|-- reviewText_clean: array (nullable = true)
|   |-- element: string (containsNull = true)
|-- summary_clean: array (nullable = true)
|   |-- element: string (containsNull = true)
```

```
[14]: df_stop.show(5)
```

```
+-----+-----+-----+-----+-----+-----+-----+
|-----+-----+
|      asin|overall| reviewTime|      reviewerID|HelpfulRecords|weightedRating|
reviewText_clean|      summary_clean|
+-----+-----+-----+-----+-----+-----+-----+
|B00J4S6YWC|      5|06 21, 2014| AUSB91MCI3WM|      0.0|
5.0|[arc, provided, a...|      [sexy, cruise]|
|B00HCZUBH8|      5| 03 3, 2014|A141H51I3H4B1S|      0.5|
5.0|[wild, ride, nanc...|[changing, gears,...|
|B006RZNR3Y|      5|07 10, 2014| AP8TKDM76TROZ|      0.0|
4.0|[well, thought, s...| [take, kindly, no!]|
|B006RZNR3Y|      3| 02 1, 2014|A22GGHISKRVAOX|      0.0|
4.0|[autistic, freque...|[mediocre, scienc...|
|B00J47H8H8|      3|03 21, 2014|A19DWIC1T7127Y|      0.75|
3.0|[book, four, five...| [losing, interest]|
+-----+-----+-----+-----+-----+-----+-----+
|-----+-----+
only showing top 5 rows
```

1 Exploratory Data Analysis

```
[15]: df_stop.describe().show()
```

```
+-----+-----+-----+-----+-----+
+-----+
|summary|      asin|      overall|reviewTime|  reviewerID|
HelpfulRecords|  weightedRating|
+-----+-----+-----+-----+-----+
+-----+
| count|      4880|      4880|      4880|      4880|
4880|      4880|
| mean|      null|4.340573770491804|      null|
null|0.3715991527158007| 4.34097108502769|
| stddev|      null|0.973934363172232|      null|
null|0.4611430329911328|0.9374090879340996|
| min|B000SRGF2W|      1|01 1, 2011| AOJVIONYIOT2|
0.0|      1.0|
| max|B00LYPZIX0|      5|12 9, 2013|AZZFSLSL2LE4FX|
1.0| 5.0000000000000001|
+-----+-----+-----+-----+-----+
+-----+
```

1.1 Word 2 Vec

```
[23]: word_vec=df_stop.select("reviewText_clean")
```

```
[32]: word_vec.show(5)
```

```
+-----+
|      summary_clean|
+-----+
|      [sexy, cruise]|
|[changing, gears,...|
|[take, kindly, no!]|
|[mediocre, scienc...|
|[losing, interest]|
+-----+
only showing top 5 rows
```

```
[27]: from pyspark.ml.feature import Word2Vec
```

```
word2Vec = Word2Vec(vectorSize=5, seed=42, inputCol="reviewText_clean",
    ↪outputCol="model")
```

```
word2Vec.setMaxIter(10)
#Word2Vec...
word2Vec.getMaxIter()
10
word2Vec.clear(word2Vec.maxIter)
model = word2Vec.fit(word_vec)
model.getMinCount()
5
model.setInputCol("words_clean")

#Word2VecModel...
model.getVectors().show(10,truncate=False)
```

```
+-----+-----+
-----+
|word      |vector
|
+-----+-----+
-----+
|clarissa  |[-0.08527789264917374,-0.061181358993053436,0.04229581356048584,0.13
291782140731812,-0.020387664437294006] |
|incident  |[-0.17075666785240173,0.026323946192860603,-0.06936486065387726,0.02
8995148837566376,-0.1795503944158554] |
|serious   |[-0.17508743703365326,0.12439662963151932,-0.06705012172460556,0.103
28210890293121,-0.18280819058418274] |
|breaks    |[-0.09439557045698166,0.07781413197517395,-0.15210963785648346,0.061
19786947965622,-0.2311510592699051] |
|forgotten |[-0.0568401962518692,0.0626450851559639,-0.001981329172849655,-0.034
09483656287193,-0.0365166962146759] |
|precious  |[-0.15493319928646088,0.09503928571939468,0.053145602345466614,0.047
84063249826431,-0.03783516213297844] |
|mario     |[-0.12475521117448807,0.09339739382266998,-0.09627118706703186,0.042
47725009918213,0.02714124508202076] |
|compliment|[0.03079369105398655,0.09589443355798721,-0.04605694115161896,0.0688
2615387439728,-0.06966786086559296] |
|lover     |[-0.09675610810518265,0.05457765609025955,-0.08897153288125992,0.104
58670556545258,-0.09521284699440002] |
|terrible  |[-0.15116223692893982,0.0013384217163547873,0.10820884257555008,0.00
8466287516057491,-0.26227179169654846] |
+-----+-----+
-----+

only showing top 10 rows
```

```
[28]: word_vec=df_stop.select("summary_clean")
```

```
[30]: word2Vec = Word2Vec(vectorSize=5, seed=42, inputCol="summary_clean",
    ↪outputCol="model")
word2Vec.setMaxIter(10)
#Word2Vec...
word2Vec.getMaxIter()
10
word2Vec.clear(word2Vec.maxIter)
model = word2Vec.fit(word_vec)
model.getMinCount()
5
model.setInputCol("words_clean")

#Word2VecModel...
model.getVectors().show(10,truncate=False)
```

```
+-----+-----+
+-----+
|word      |vector
|
+-----+-----+
+-----+
|ideas     |[0.07896555960178375,0.08228708803653717,-0.0314699150621891,-0.05050
545558333397,-0.01455814577639103]
|sweet     |[-0.0714794397354126,-0.07245013117790222,-0.008758701384067535,-0.05
772051960229874,0.1064988225698471]
|beautiful |[0.016373470425605774,-0.09401625394821167,-0.09978445619344711,0.012
30132207274437,0.010742578655481339]
|writing   |[0.1669931560754776,-0.09632845968008041,-0.023153474554419518,-0.050
83288624882698,0.07983992248773575]
|funny     |[0.05766937509179115,-0.10599081218242645,-0.0632324367761612,-0.0160
7191003859043,-0.0852198451757431]
|weird     |[-0.019863391295075417,0.006063351407647133,-0.06246356666088104,-0.0
07803264074027538,0.030566997826099396]
|wow       |[-0.014897726476192474,0.053794026374816895,-0.10401398688554764,0.05
399356409907341,-0.034788161516189575]
|series,   |[0.01268547773361206,-0.08342977613210678,0.0755157321691513,0.066184
96030569077,0.0459398478269577]
|wanting   |[-0.07497256249189377,0.08413670212030411,-0.018536822870373726,0.008
434544317424297,0.06759133189916611]
|please    |[0.0951368510723114,0.02940688654780388,-0.0927429273724556,0.0893965
2889966965,0.0991884097456932]
+-----+-----+
+-----+
only showing top 10 rows
```

1.2 Word Cloud

```
[35]: #!pip install wordcloud
```

```
[101]: from wordcloud import WordCloud, ImageColorGenerator
from PIL import Image
import matplotlib.pyplot as plt
from os import path
```


```
[102]: reviews=df_stop.select("reviewText_clean")
reviews=reviews.toPandas()
```

```
[104]: text=reviews['reviewText_clean'].apply(', '.join)
```

```
[105]: text = ", ".join(review for review in text)
```

```
[106]: wordcloud = WordCloud(max_font_size=50, max_words=100,
    ↪background_color="white").generate(text)

# Display the generated image:
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")
plt.show()
```



```
[109]: wordcloud.to_file("amazon.png")
```


```
[109]: <wordcloud.wordcloud.WordCloud at 0x7f34f7095128>
```

```
[111]: reviews=df_stop.select("summary_clean")
reviews=reviews.toPandas()
```

```
text=reviews['summary_clean'].apply(', '.join)
text = ", ".join(review for review in text)
```

```
[112]: wordcloud = WordCloud(max_font_size=50, max_words=100,
    ↪background_color="white").generate(text)

# Display the generated image:
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")
plt.show()
```



```
[113]: wordcloud.to_file("summary.png")
```


```
[113]: <wordcloud.wordcloud.WordCloud at 0x7f34f6a7a2e8>
```

1.3 Histogram


```
[117]: df_pd=df_stop.toPandas()
```

```
[134]: df_pd.overall.value_counts().plot(kind='bar')
plt.xlabel('Rating')
plt.ylabel('Number of Reviews')
plt.title("Count of Ratings")
```

```
[134]: Text(0.5, 1.0, 'Count of Ratings')
```

```
[138]: df_pd.reviewerID.value_counts().head(10).plot(kind = 'bar')
plt.xticks(rotation = 70)
plt.xlabel('UserID')
plt.ylabel('Number of Reviews')
plt.title("Top Reviewers By Number of Reviews")
plt.show()
```



[]: